DEEP LEARNING INSTITUTE

不用写代码
使用NVIDIA DIGITS进行图像分类

侯宇涛
Developer marketing Director
Certified Instructor, NVIDIA Deep Learning Institute
NVIDIA China

ISAAC VIDEO

https://www.youtube.com/watch?v=PIDJdKwo2VM

NVIDIA Isaac SDK

Accelerate Your Creation of Autonomous Machines

The NVIDIA Isaac software development kit (SDK) makes it easy for developers to create and deploy AI-powered robotics. The SDK is a collection of libraries, drivers, APIs, and other tools that will save you hundreds of hours by making it easy to add AI into next-generation robots for perception, navigation, and manipulation.

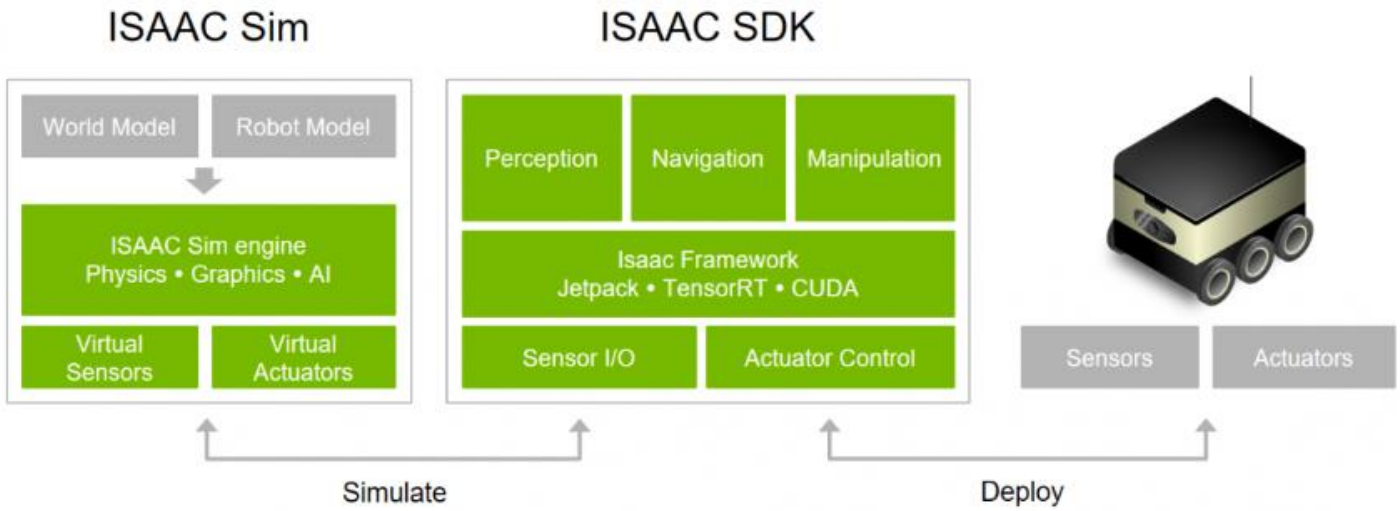

Benefits

- Speed up your development cycle using a robust robotics framework
- Train and test your robot in simulation and transfer it to physical robots
- Solve your robotics challenges using high-performance GPU-accelerated algorithms

动手实验培训

通用基础 +专业应用领域

Caffe2

Microsoft Cognitive Toolkit

mxnet

TensorFlow

PYTORCH

Fundamentals

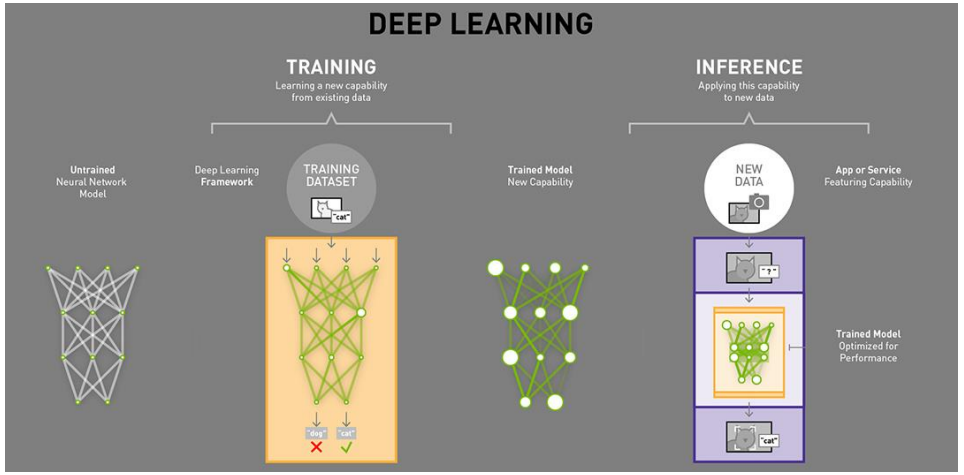

Autonomous Vehicles

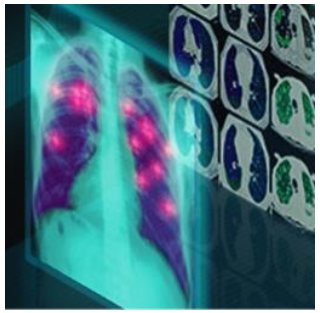

Healthcare

Game Development and Digital Content

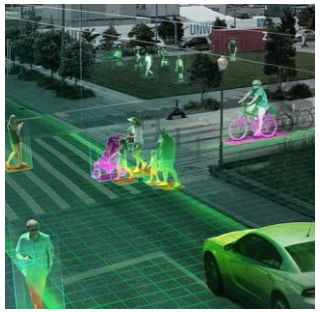

Intelligent Video Analytics

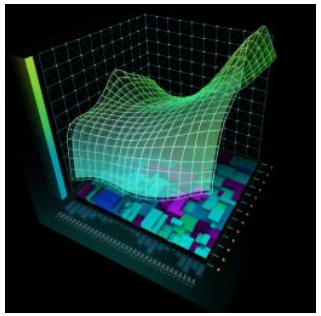

Finance

More content in a range of topics coming soon…

NVIDIA DEEP LEARNING INSTITUTE

Traditional Deep Learning Workflow

- DL frameworks, Caffe, etc. aimed at computer scientist not data scientist
- Juggle multiple files & windows
- Handcrafted visualizations
- Manual log file parsing
- Manual experiment logging
- Model editing in Lua IDE files

```
111     -- stage 1 : filter bank -> squashing -> L2 pooling -> normalization
112     model:add(nn.SpatialConvolutionMM(nfeats, nstates[1], filtsize, filtsize))
113     model:add(nn.Tanh())
114     model:add(nn.SpatialLPPooling(nstates[1],2,poolsize,poolsize,poolsize,poolsize))
115     model:add(nn.SpatialSubtractiveNormalization(nstates[1], normkernel))
116
117     -- stage 2 : filter bank -> squashing -> L2 pooling -> normalization
118     model:add(nn.SpatialConvolutionMM(nstates[1], nstates[2], filtsize, filtsize))
119     model:add(nn.Tanh())
120     model:add(nn.SpatialLPPooling(nstates[2],2,poolsize,poolsize,poolsize,poolsize))
121     model:add(nn.SpatialSubtractiveNormalization(nstates[2], normkernel))
122
123     -- stage 3 : standard 2-layer neural network
124     model:add(nn.Reshape(nstates[2]*filtsize*filtsize))
125     model:add(nn.Linear(nstates[2]*filtsize*filtsize, nstates[3]))
126     model:add(nn.Tanh())
127     model:add(nn.Linear(nstates[3], noutputs))
128   end
129 else
130
131   error('unknown -model')
132
133 end
134
135 ----------------------------------------------------------------------
136 print '==> here is the model:'
137 print(model)
```

```
==> here is the model:
nn.Sequential {
  [input -> (1) -> (2) -> (3) -> (4) -> (5) -> (6) -> (7) -
  (1): nn.SpatialConvolutionMM(3 -> 64, 5x5)
  (2): nn.Tanh
  (3): nn.Sequential {
    [input -> (1) -> (2) -> (3) -> (4) -> output]
    (1): nn.Square
    (2): nn.SpatialAveragePooling(2,2,2,2)
    (3): nn.MulConstant
    (4): nn.Sqrt
  }
  (4): nn.SpatialSubtractiveNormalization
  (5): nn.SpatialConvolutionMM(64 -> 64, 5x5)
  (6): nn.Tanh
  (7): nn.Sequential {
    [input -> (1) -> (2) -> (3) -> (4) -> output]
    (1): nn.Square
    (2): nn.SpatialAveragePooling(2,2,2,2)
    (3): nn.MulConstant
    (4): nn.Sqrt
  }
  (8): nn.SpatialSubtractiveNormalization
  (9): nn.Reshape(1600)
  (10): nn.Linear(1600 -> 128)
  (11): nn.Tanh
  (12): nn.Linear(128 -> 10)
}
```

```
306 #28p        a02 Take the trueNLL model w/ kappa_true=0.76 from expt#20  -> true + MSE under settings
307                     d   all data,       55 epochs,  epochSize 500,  regime with LR=0.2
308 th main_custom_noTest.lua -cache devAnnealAugFinal -data ../data-1024-fulltrain -nGPU 4 -name TrueNLL
    nDonkeys 6 -nEpochs 55 -epochSize 500 -batchSize 32 -LR 0.2 -netType 1024/nin-preall-multi -retrain /
    multiGPU/highres/learn_1024_anneal/devAnnealAug/EvenNLLtoTrueNLL_fixedLR,CRfinal=trueCR,LR=0.1,batchS
    nGPU=4,netType=1024/nin-preall-multi-fixedLR/05-48_on_25-Jul-15/model_100.t7 -CRinitial trueCR -CRfin
    1 -manualSeed 4359 | tee expt28p_ap02_TrueNLLtoTrueMSEfinal.log
309
310
311 #29p        a03 Take the 62411+0.75MSE model  w/ kappa_true=0.78 from expt#24 -> true + MSE under set
312                     b   regular data,    55 epochs,  epochSize 200,  regime with LR=0.2
313 th main_custom.lua -cache devAnnealAug -data ../data-1024 -nGPU 4 -name CR62411MSE3by4_seed864521_exp
    epochSize 200 -batchSize 32 -LR 0.2 -netType 1024/nin-preall-multi -retrain /home/mnt/retinopathy/ima
    learn_1024_anneal/devAnnealAug/pretrained62411MSE3by4_trainFurther,CRfinal=6,2,4,1,1,CRinitial=6,2,4,
    MSEinitial=0.75,batchSize=32,epochSize=300,nDonkeys=6,nGPU=4,netType=1024/nin-preall-multi-fixedLR/05
    CRinitial 6,2,4,1,1 -CRfinal trueCR -MSEinitial 0 -MSEfinal 1 -manualSeed 864521| tee expt29p_a03_CR
314
315
316 #30p        ap06    Take the 62411+0.75MSE model  w/ kappa_true=0.78 from expt#24 -> true + MSE under
317                     d   all data,       55 epochs,  epochSize 500,  regime with LR=0.2
318 th main_custom_noTest.lua -cache devAnnealAugFinal -data ../data-1024-fulltrain -nGPU 4 -name CR62411
    nDonkeys 6 -nEpochs 55 -epochSize 500 -batchSize 32 -LR 0.2 -netType 1024/nin-preall-multi -retrain /
    multiGPU/highres/learn_1024_anneal/devAnnealAug/pretrained62411MSE3by4_trainFurther,CRfinal=6,2,4,1,1
    MSEfinal=0.75,MSEinitial=0.75,batchSize=32,epochSize=300,nDonkeys=6,nGPU=4,netType=1024/nin-preall-mu
    model_55.t7 -CRinitial 6,2,4,1,1 -CRfinal trueCR -MSEinitial 0 -MSEfinal 1 -manualSeed 351987 | tee e
    log
319
```

课程介绍：

级别：初级 | 预备知识：无
行业：所有 | Frameworks: Caffe

此实验室会向您展示如何通过在 Caffe 框架上的 NVIDIA DIGITS 和MNIST 手写数据集，在深度学习工作流程中利用深度神经网络 (DNN)， 尤其是卷积神经网络 (CNN) 解决真实图像分类问题，您会学到：

1. 构建运行在GPU上的深度神经网络
2. 管理数据准备、模型定义、模型训练和问题排查过程
3. 使用验证数据来测试和尝试不同策略来提升模型性能

完成此实验室后，您将能够使用 NVIDIA DIGITS 来构建、训练、评估和提升您的图像分类应用程序中 CNN 的准确性。

NVIDIA DIGITS

交互式深度学习GPU训练系统

Process Data
加载数据集

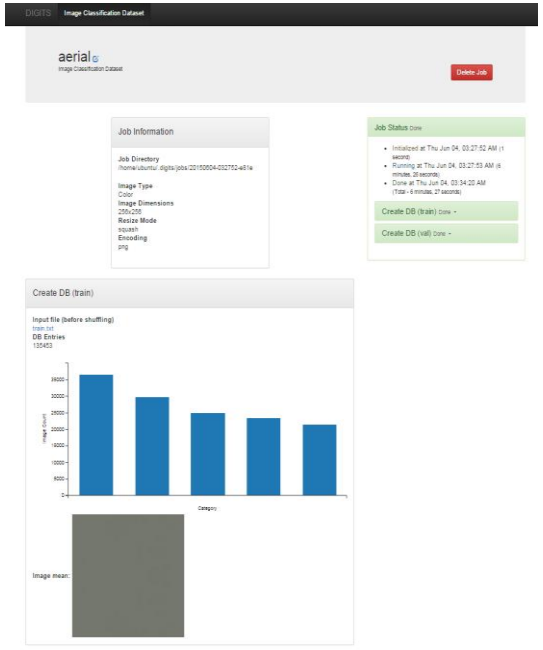

Configure DNN
配置神经网络

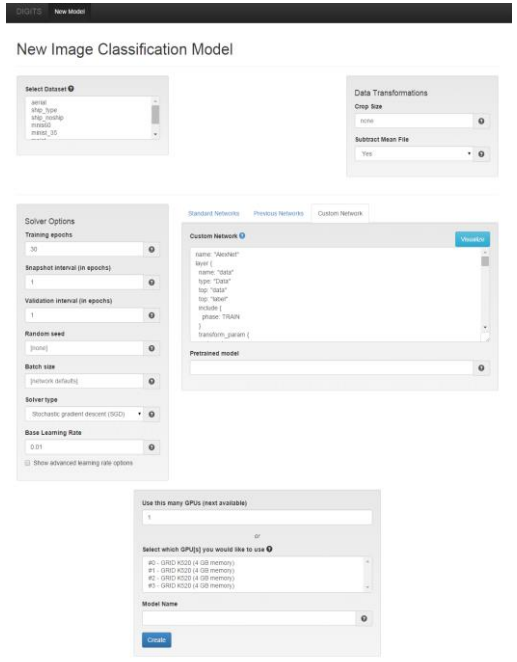

Monitor Progress
监控训练过程

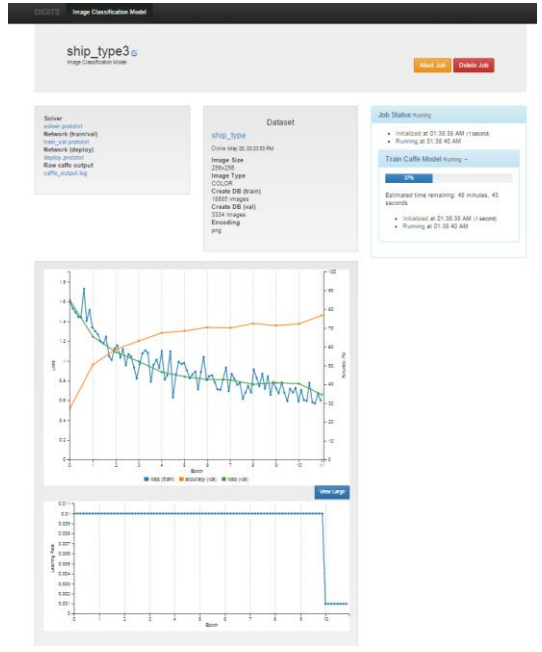

Visualization
可视化校验

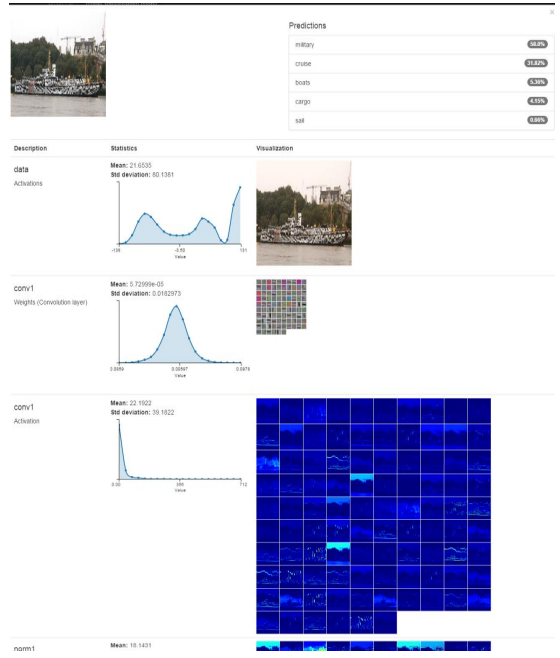

NVIDIA’S DIGITS

交互式深度学习GPU训练系统

- 简化通用的深度学习任务如：
  - 管理数据
  - 在多GPU系统上设计并训练神经网络
  - 使用高级可视化界面，监控实时性能
- 完全的交互式界面，使得数据科学家可以专注在设计及训练网络，节约编程及调试代码的时间
- 开源

DIGITS - HOME

Clicking DIGITS will bring you to this Home screen

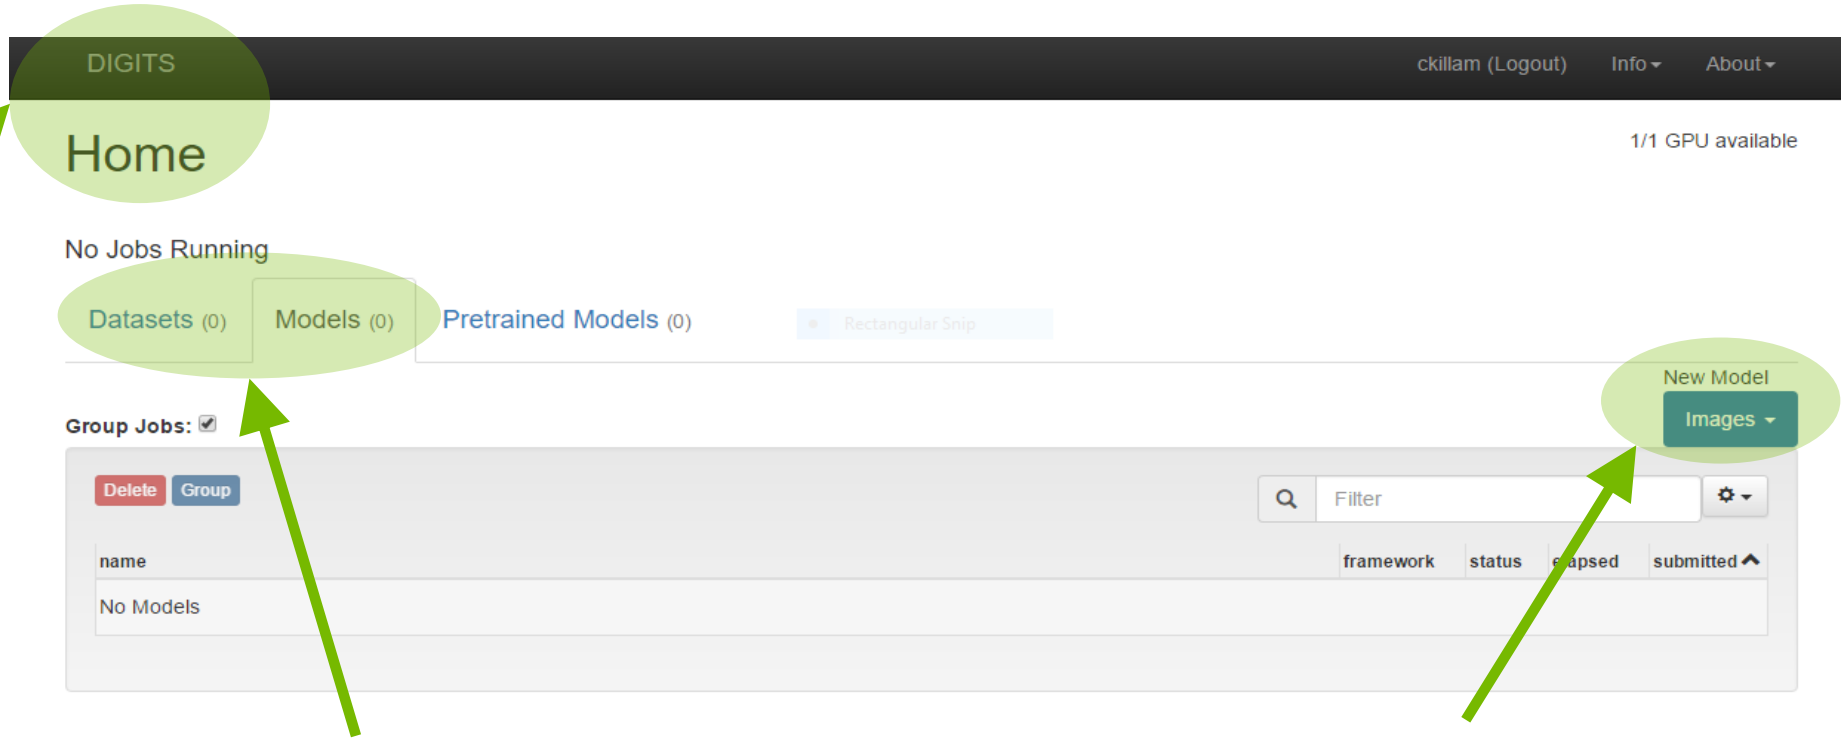

Click here to see a list of existing datasets or models

Clicking here will present different options for model and dataset creation

DIGITS - DATASET

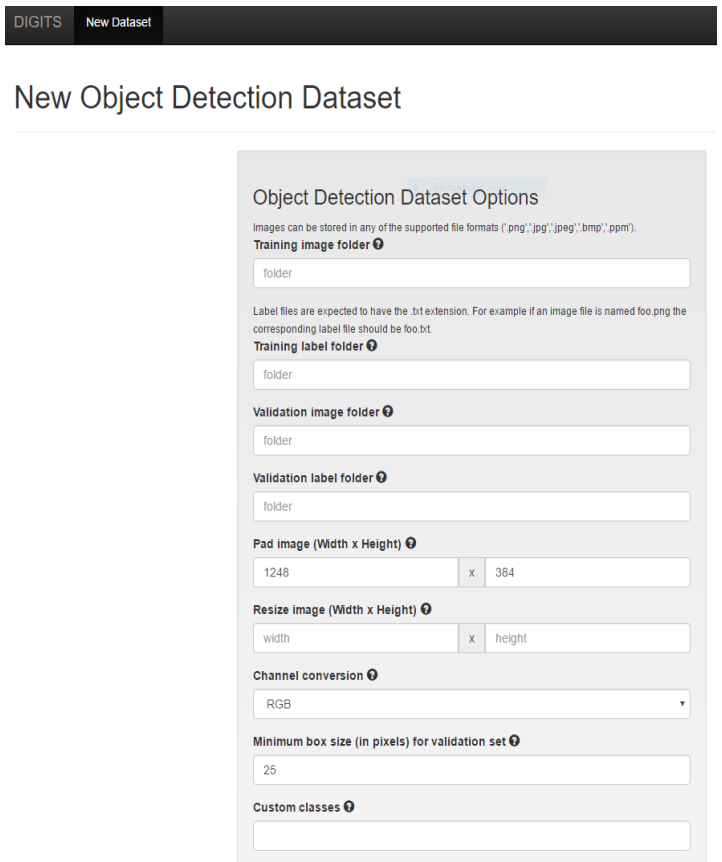

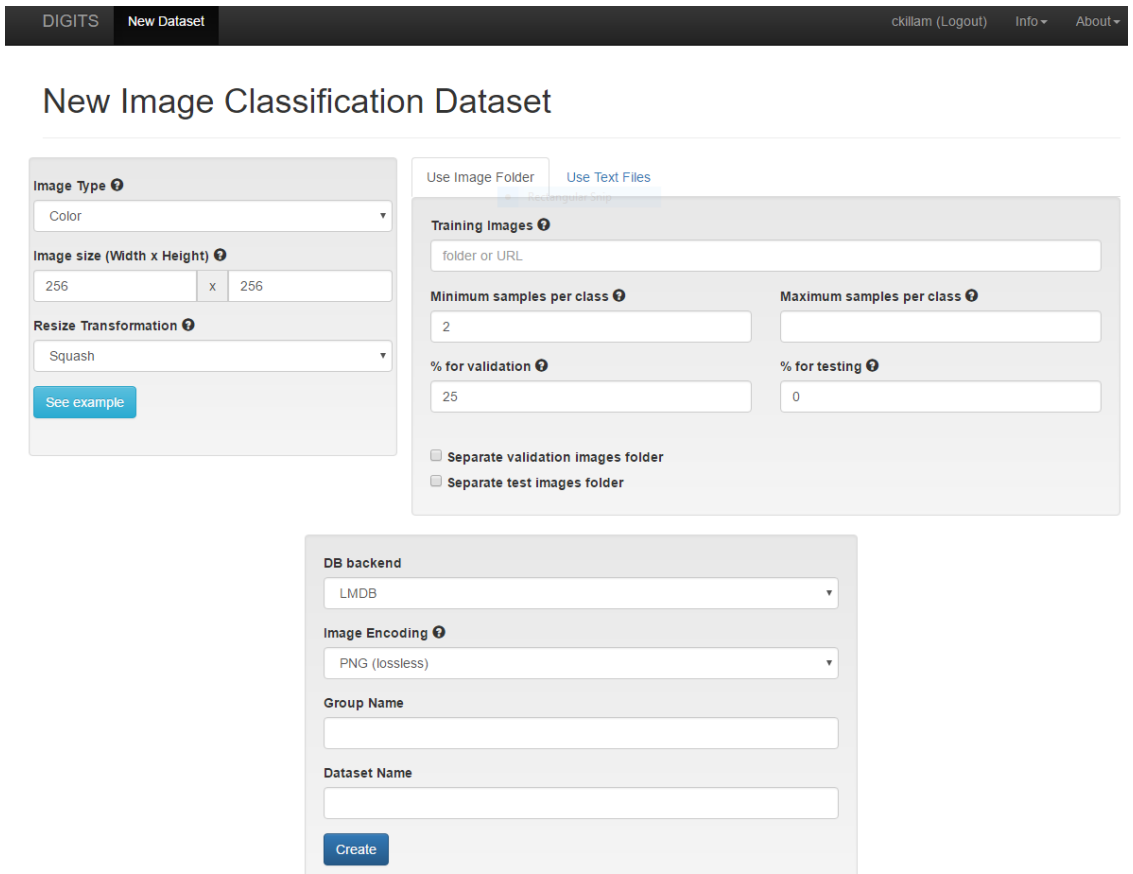

Different options will be presented based upon the task

DIGITS - MODEL

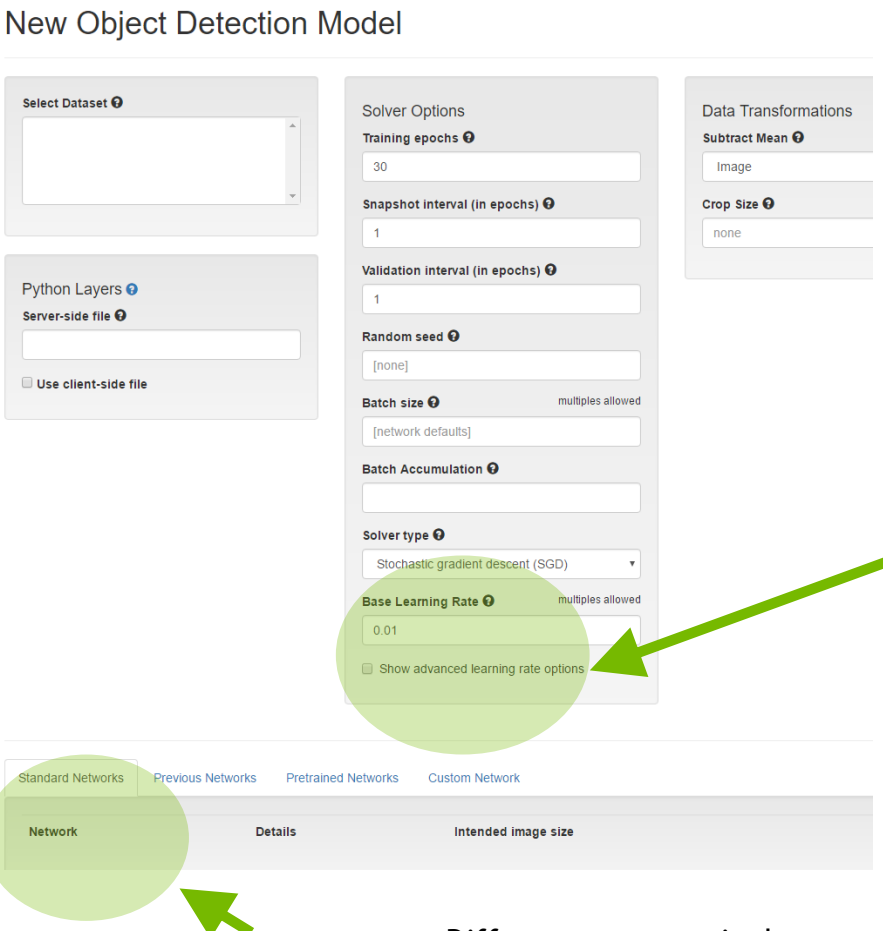

Define custom layers with Python

Can anneal the learning rate

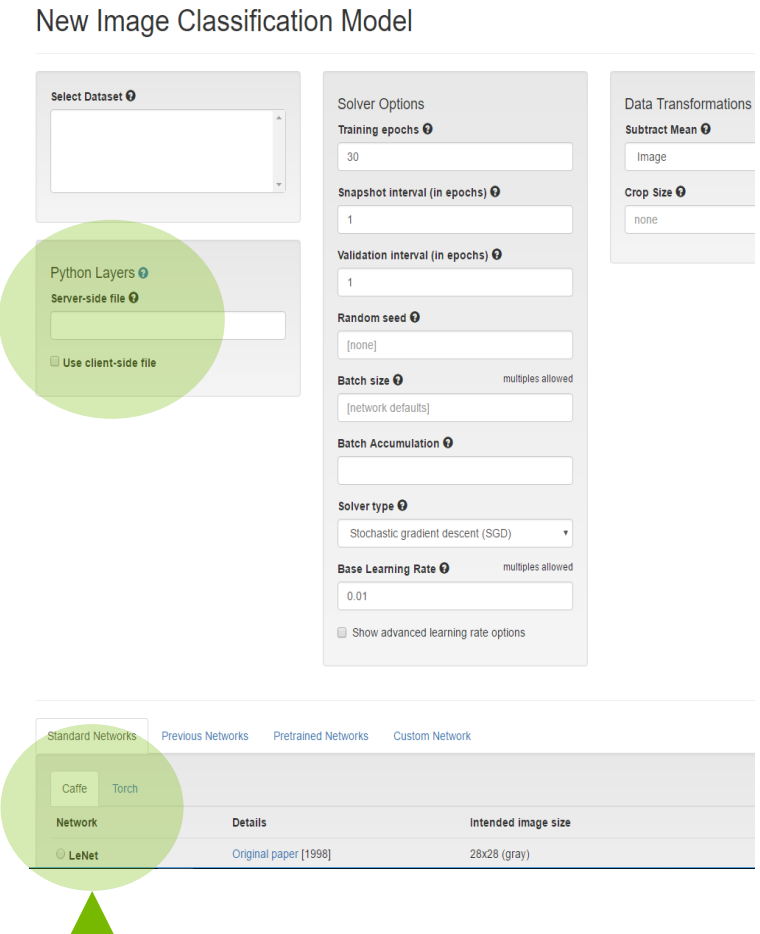

Differences may exist between model tasks

DIGITS

- MODEL

DIGITS | New Model

Standard Networks | Previous Networks | Pretrained Networks | Custom Network

Caffe | Tensorflow

Custom Network Visualize

```
  top: "scaled"
  type: "Power"
  power_param {
    # 1/(standard deviation on MNIST dataset)
    scale: 0.0125
  }
}
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "scaled"
  top: "conv1"
  param {
    lr_mult: 1
  }
  param {
    lr_mult: 2
  }
  convolution_param {
    num_output: 20
    kernel_size: 5
    stride: 1
    weight_filler {
      type: "xavier"
    }
    bias_filler {
      type: "constant"
    }
  }
}
layer {
  name: "pool1"
  type: "Pooling"
  bottom: "conv1"
  top: "pool1"
  pooling_param {
    pool: MAX
    kernel_size: 2
    stride: 2
  }
}
layer {
  name: "conv2"
```

Pretrained model(s)

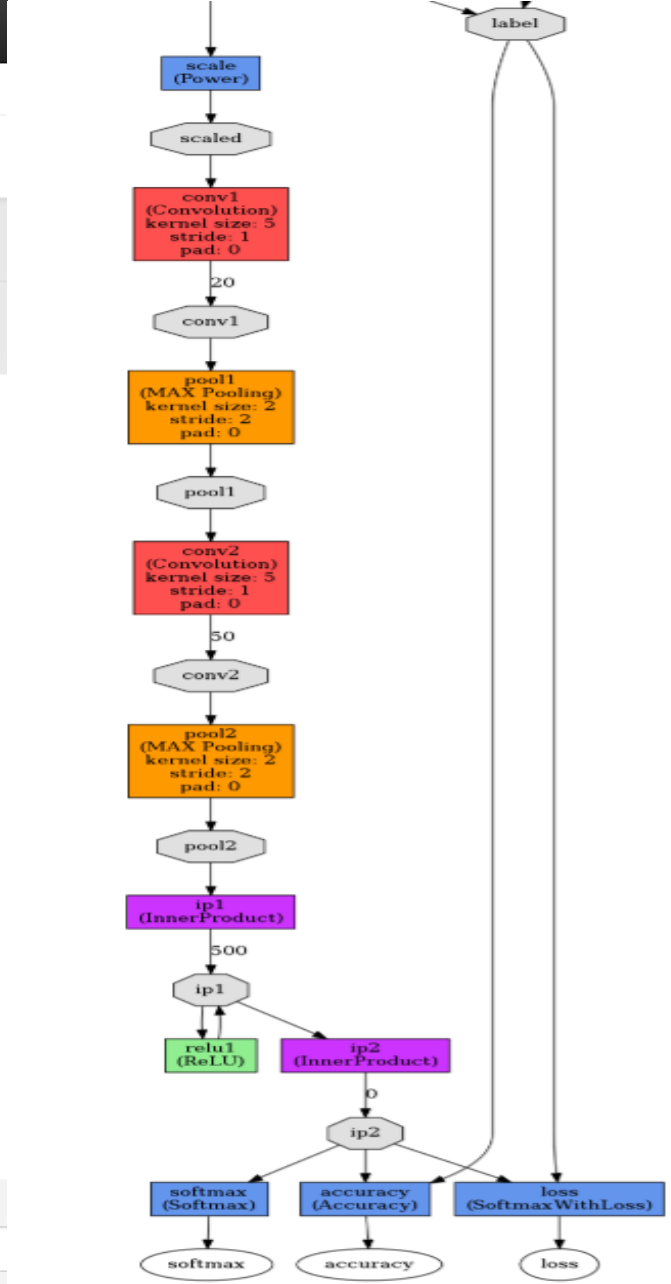

DIGITS - TRAINING

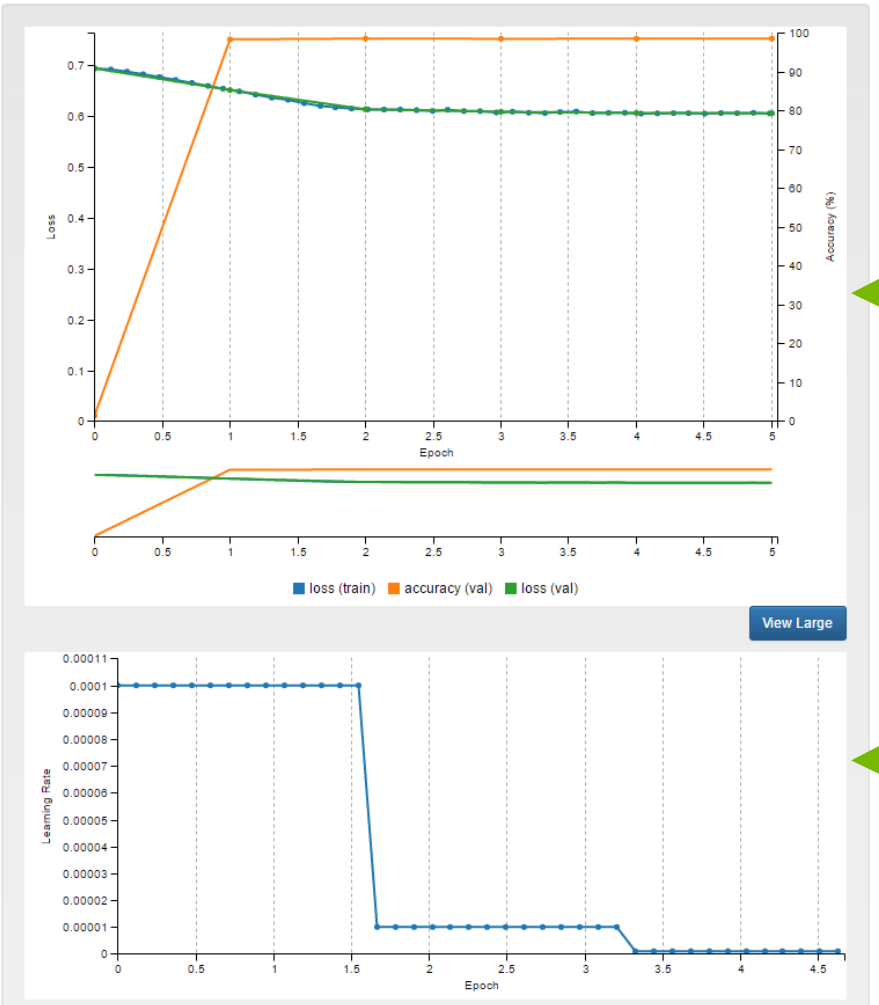

Loss function and accuracy during training

Annealed learning rate

NVIDIA DIGITS

Visualization

Predictions

military	58.0%
cruise	31.82%
boats	5.36%
cargo	4.15%
sail	0.66%

Description	Statistics	Visualization
data Activations	**Mean:** 21.6535 **Std deviation:** 80.1381	
conv1 Weights (Convolution layer)	**Mean:** 5.72999e-05 **Std deviation:** 0.0182973	
conv1 Activation	**Mean:** 22.1922 **Std deviation:** 39.1822	
norm1	**Mean:** 18.1431	

NVIDIA DIGITS

Trained Models

Select Model

Epoch #5 ▾ | Download Model | Make Pretrained Model | Publish to inference server

Test a single image

Image Path

Upload image

Browse...

Show visualizations and statistics

Classify One

Test a list of images

Upload Image List

Browse...

Accepts a list of filenames or urls (you can use your val.txt file)

Image folder *(optional)*

Relative paths in the text file will be prepended with this value before reading

Number of images use from the file

All

Leave blank to use all

Classify Many

Number of images to show per category

9

Top N Predictions per Category

20180708-064929-6e4b_epoch_10.0.tar.gz 7/8/2018 6:50 AM WinRAR 压缩文件 1,561 KB

名称	大小	压缩后大小	类型
..			Local Disk
deploy.prototxt	1,823	?	PROTOTXT File
info.json	769	?	JSON File
labels.txt	20	?	Text Document
mean.binaryproto	3,147	?	BINARYPROTO Fil
original.prototxt	2,346	?	PROTOTXT File
snapshot_iter_7040.caffemodel	1,725,144	?	CAFFEMODEL File
solver.prototxt	300	?	PROTOTXT File
train_val.prototxt	2,600	?	PROTOTXT File

DIGITS DEEP LEARNING Workflows

IMAGE CLASSIFICATION

98% Dog

2% Cat

Classify images into classes or categories

Object of interest could be anywhere in the image

OBJECT DETECTION

Find instances of objects in an image

Objects are identified with bounding boxes

IMAGE SEGMENTATION

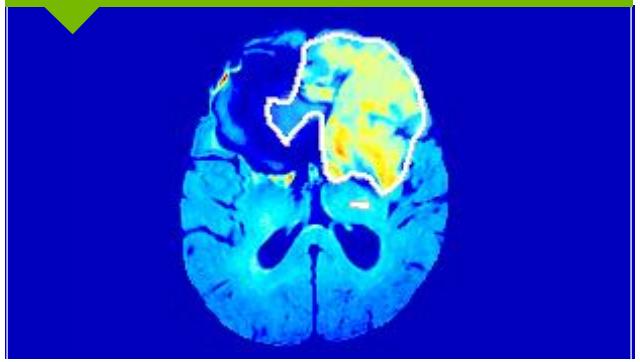

Partition image into multiple regions

Regions are classified at the pixel level

NVIDIA digits 6

TENSORFLOW SUPPORT

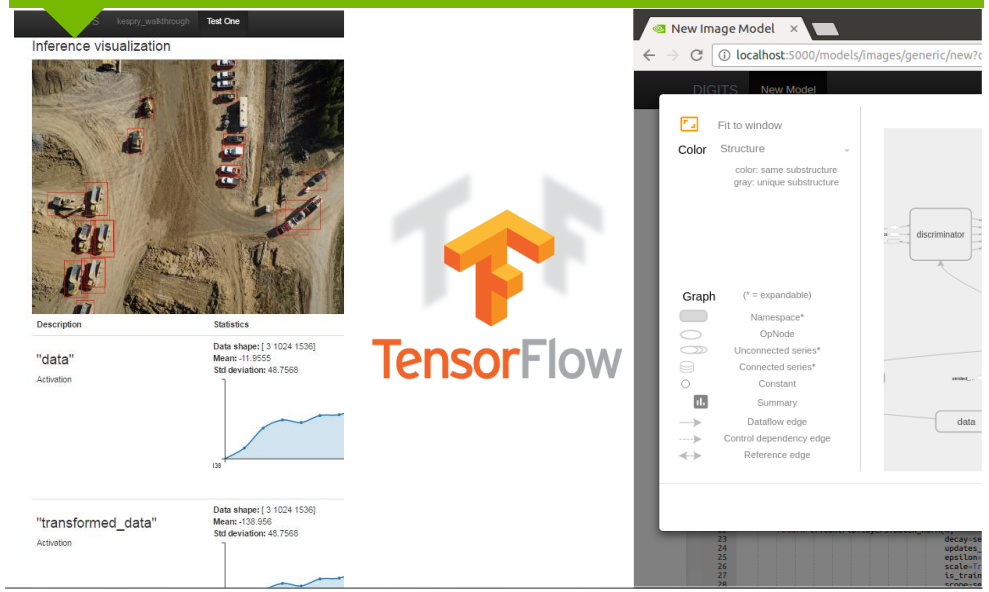

Train TensorFlow Models Interactively with DIGITS

NEW PRE-TRAINED MODELS

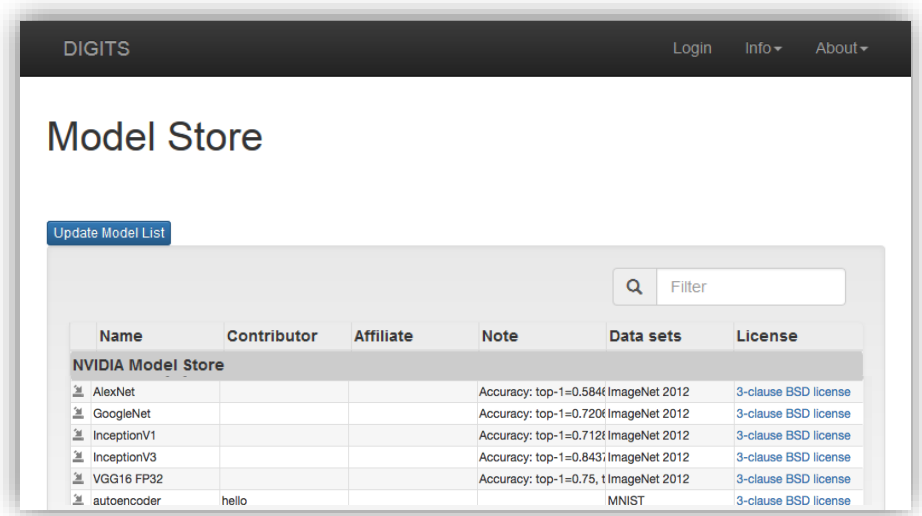

Image Classification: VGG-16, ResNet50
Object Detection: DetectNet

DIGITS 6 is now available now as a free download to members of NVIDIA Developer Program

developer.nvidia.com/digits

TUTORIAL教程:

HANDWRITTEN DIGIT RECOGNITION
手写体数字识别模型训练及优化

HANDWRITTEN DIGIT RECOGNITION

HELLO WORLD of machine learning?

MNIST data set of handwritten digits from Yann Lecun's website

All images are 28x28 grayscale

Pixel values from 0 to 255

60k training examples, 10k test examples

Input vector of size 784

Output value is integer from 0-9

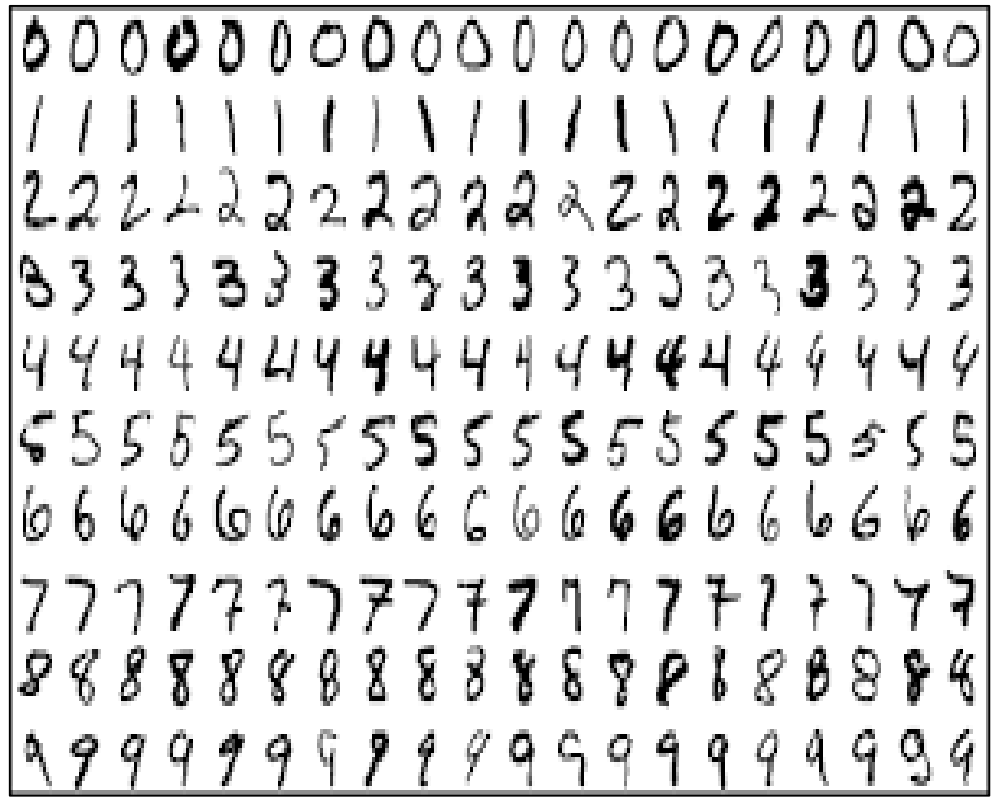

SMALL DATASET

6000 x images

- Dataset
  - Training Images: /home/ubuntu/data/train_small
- Model
  - “MNIST small”

FIRST RESULTS

Small dataset (10 epochs)

- 96% of accuracy achieved
- Training is done within one minute

	SMALL DATASET
	1 : 99.90 %
	2 : 69.03 %
	8 : 71.37 %
	8 : 85.07 %
	0 : 99.00 %
	8 : 99.69 %
	8 : 54.75 %

FULL DATASET

6x larger dataset

- Dataset
  - Training Images: /home/ubuntu/data/train_full
  - Dataset Name: MNIST full
- Model
  - Clone “MNIST small”.
  - Give a new name “MNIST full” to push the create button

SECOND RESULTS

Full dataset (10 epochs)

- 99% of accuracy achieved
- No improvements in recognizing real-world images

	SMALL DATASET	FULL DATASET
	1 : 99.90 %	0 : 93.11 %
	2 : 69.03 %	2 : 87.23 %
	8 : 71.37 %	8 : 71.60 %
	8 : 85.07 %	8 : 79.72 %
	0 : 99.00 %	0 : 95.82 %
	8 : 99.69 %	8 : 100.0 %
	8 : 54.75 %	2 : 70.57 %

DATA AUGMENTATION

Adding Inverted Images

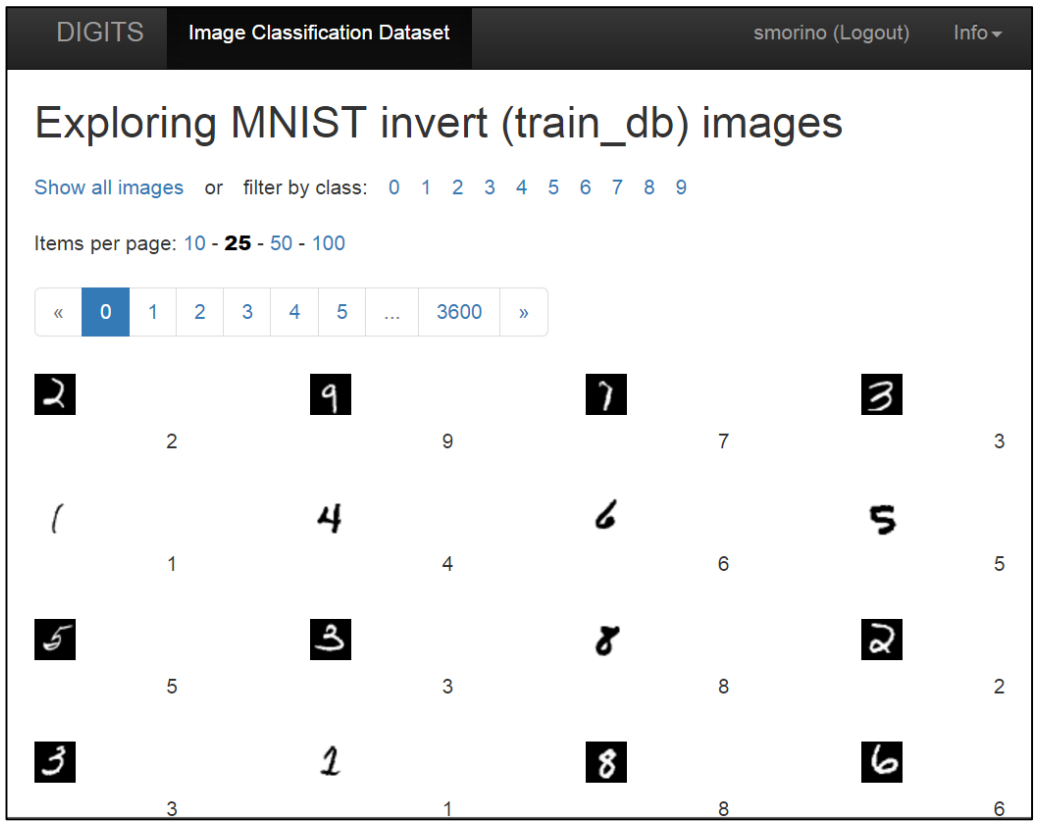

- Pixel(Inverted) = 255 – Pixel(original)
- White letter with black background
  - Black letter with white background
- Training Images: /home/ubuntu/data/train_invert
- Dataset Name: MNIST invert

DATA AUGMENTATION

Adding inverted images (10 epochs)

	SMALL DATASET	FULL DATASET	+INVERTED
	1 : 99.90 %	0 : 93.11 %	1 : 90.84 %
	2 : 69.03 %	2 : 87.23 %	2 : 89.44 %
	8 : 71.37 %	8 : 71.60 %	3 : 100.0 %
	8 : 85.07 %	8 : 79.72 %	4 : 100.0 %
	0 : 99.00 %	0 : 95.82 %	7 : 82.84 %
	8 : 99.69 %	8 : 100.0 %	8 : 100.0 %
	8 : 54.75 %	2 : 70.57 %	2 : 96.27 %

MODIFY THE NETWORK

Adding filters and ReLU layer

```
layer {
        name: "pool1“
        type: "Pooling“
        …
}

layer {
        name: "reluP1"
        type: "ReLU"
        bottom: "pool1"
        top: "pool1"
}

layer {
        name: "reluP1“
```

```
layer {
  name: "conv1"
  type: "Convolution"
        ...
        convolution_param {
        num_output: 75
        ...
layer {
        name: "conv2"
        type: "Convolution"
        ...
        convolution_param {
        num_output: 100
        ...
```

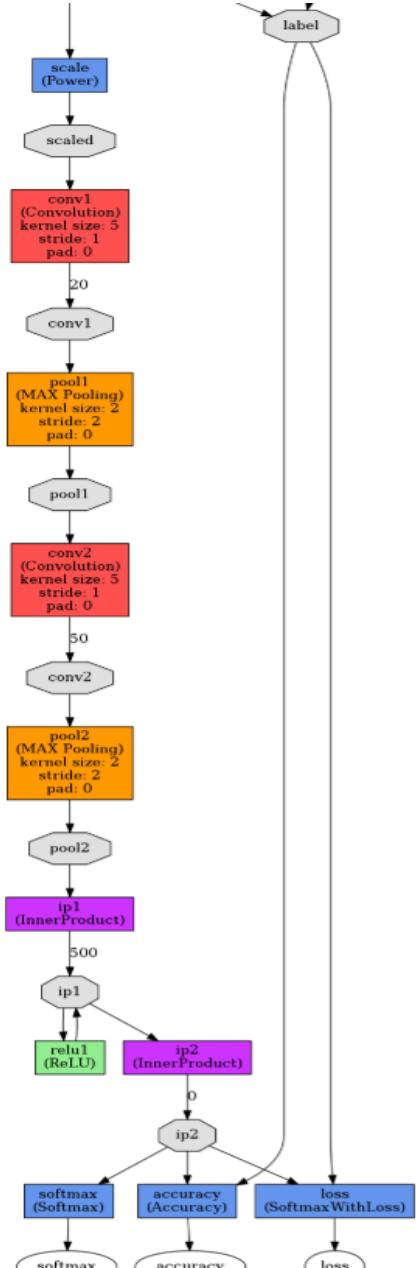

MODIFY THE NETWORK

Adding ReLU Layer

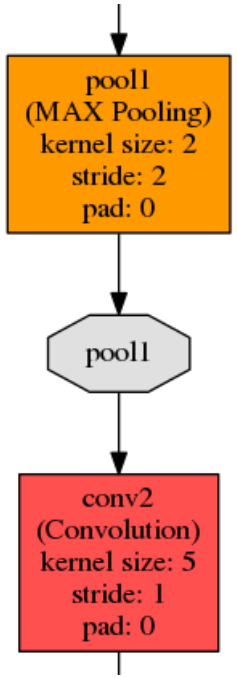

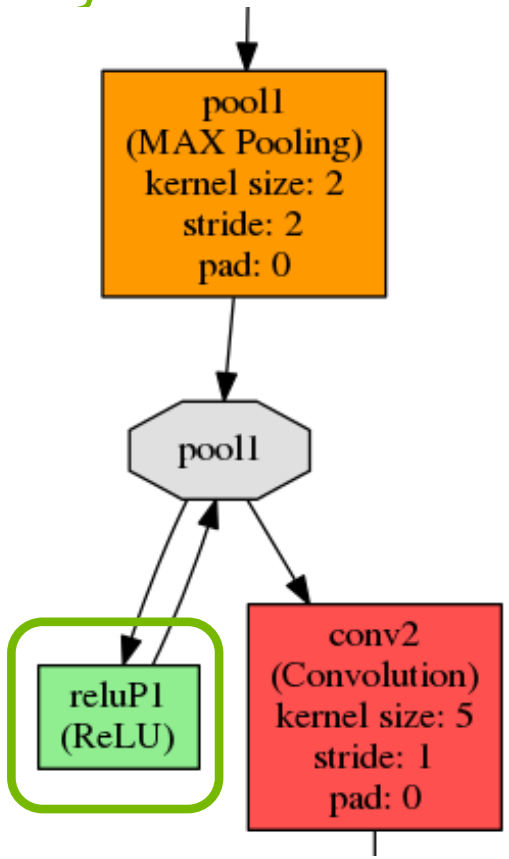

MODIFIED NETWORK

Adding filters and ReLU layer (10 epochs)

	SMALL DATASET	FULL DATASET	+INVERTED	ADDING LAYER
	1 : 99.90 %	0 : 93.11 %	1 : 90.84 %	1 : 59.18 %
	2 : 69.03 %	2 : 87.23 %	2 : 89.44 %	2 : 93.39 %
	8 : 71.37 %	8 : 71.60 %	3 : 100.0 %	3 : 100.0 %
	8 : 85.07 %	8 : 79.72 %	4 : 100.0 %	4 : 100.0 %
	0 : 99.00 %	0 : 95.82 %	7 : 82.84 %	2 : 62.52 %
	8 : 99.69 %	8 : 100.0 %	8 : 100.0 %	8 : 100.0 %
	8 : 54.75 %	2 : 70.57 %	2 : 96.27 %	8 : 70.83 %

NVIDIA DIGITS

Trained Models

Select Model

Epoch #5 ▾ | Download Model | Make Pretrained Model | Publish to inference server

Test a single image

Image Path

Upload image

Browse...

Show visualizations and statistics

Classify One

Test a list of images

Upload Image List

Browse...

Accepts a list of filenames or urls (you can use your val.txt file)

Image folder *(optional)*

Relative paths in the text file will be prepended with this value before reading

Number of images use from the file

All

Leave blank to use all

Classify Many

Number of images to show per category

9

Top N Predictions per Category

20180708-064929-6e4b_epoch_10.0.tar.gz 7/8/2018 6:50 AM WinRAR 压缩文件 1,561 KB

名称	大小	压缩后大小	类型
..			Local Disk
deploy.prototxt	1,823	?	PROTOTXT File
info.json	769	?	JSON File
labels.txt	20	?	Text Document
mean.binaryproto	3,147	?	BINARYPROTO Fil
original.prototxt	2,346	?	PROTOTXT File
snapshot_iter_7040.caffemodel	1,725,144	?	CAFFEMODEL File
solver.prototxt	300	?	PROTOTXT File
train_val.prototxt	2,600	?	PROTOTXT File

How to get DIGITS

Two methods to install DIGITS

Simple way:

- OS – Ubuntu14.04
- Download link : https://developer.nvidia.com/digits

Others (from source code):

- OSX，Windows (not tested)
- Download NVIDIA-Caffe: https://github.com/NVIDIA/caffe
- Download Digits: https://github.com/NVIDIA/DIGITS

Recommended HW/SW environment:

- GPU Compute Capability > 3.0 (Kepler and later)，cuDNN v5
- OS – Ubuntu14.04

NVIDIA 深度学习学院（DLI）

深度学习和加速计算培训
面向开发者、数据科学家和研究人员

- 真 · 实战培训，云端完全配置的 GPU 实验环境
- 全球同步，与 TensorFlow facebook 等机构共创课程
- 系统化培训，提升解决行业实际问题的能力
- NVIDIA “开发者认证”证书，构建职场竞争力

DLI 课程和开发者资源
“NVIDIA 开发者社区”

在线免费实训课
在 GPU 云上训练图像分类模型

不用写代码
使用NVIDIA DIGITS进行图像分类

7月26日 20:00-22:00
扫码进群